

# Inference and Modeling Intro

Spring R '24

Dominic Bordelon, Research Data Librarian, ULS

# Agenda

1. Inference and hypothesis testing
  1. Distribution functions
  2. Hypothesis tests
2. Modeling (simple linear regression)

# About the trainer

**Dominic Bordelon, Research Data Librarian**

University Library System, University of Pittsburgh

[dbordelon@pitt.edu](mailto:dbordelon@pitt.edu)

Services for the Pitt community:

- Consultations
- Training (on-request and via public workshops)
- Talks (on-request and publicly)
- Research collaboration

Support areas and interests:

- Computer programming fundamentals, esp. for data processing and analysis
- Open Science and Data Sharing
- Data stewardship/curation
- Research methods; science and technology studies

# Today's packages

**stats** (part of base R, automatically attaches)

**infer**: tidy inference

**tidymodels**, particularly:

- **broom**: tidy model representation
- **parsnip**: standardized modeling interface

```
1 install.packages(c("tidymodels", "infer"))
```



# ...and using penguins examples

```
1 install.packages("palmerpenguins")
```

```
1 library(palmerpenguins)
```

```
2
```

```
3 # load palmerpenguins' data into your environment:
```

```
4 data(penguins)
```

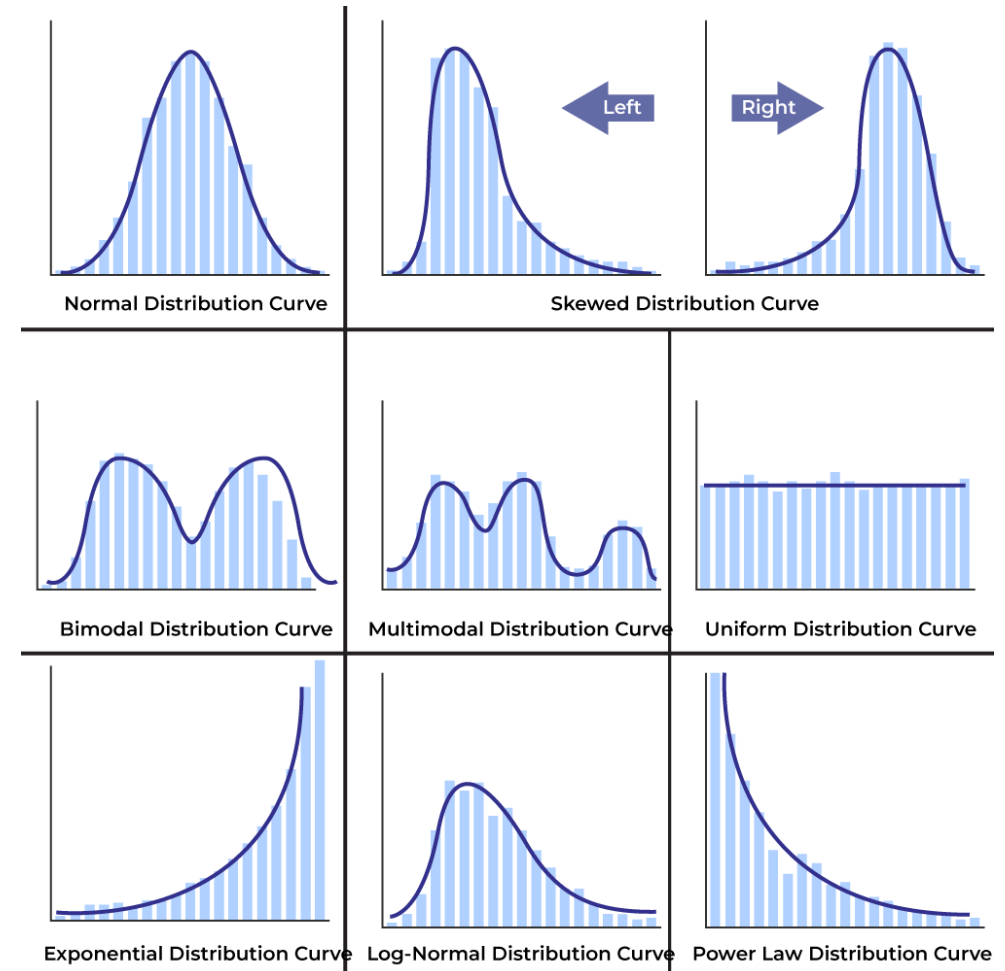
```
5 names(penguins)
```

```
[1] "species"          "island"           "bill_length_mm"  
[4] "bill_depth_mm"    "flipper_length_mm" "body_mass_g"  
[7] "sex"              "year"
```

# Inference and hypothesis testing


# Frequency distributions

- In statistics we want to estimate parameters of the *population* using observed *samples*
- In frequentist or parametric statistics, we assume that the population distribution can be approximated with standard forms such as the Gaussian (normal) distribution



A variety of distributions. Image source: [Geeks for Geeks](#)

# Distribution functions



Probability distributions in R are typically described by four functions. The examples here are for the Normal distribution.

All have `mean` and `sd` arguments.

- `dnorm(x)` for density (height of curve) at  $x$
- `pnorm(q)` for getting a probability  $p$  value at some  $q$ , i.e.,  $P(X \leq q)$
- `qnorm(p)` for finding the quantile of some  $p$
- `rnorm(n)` for generating a random sample of size  $n$  from  $N(\mu, \sigma)$



# Distributions

Commonly used distributions in the `stats` package

Distribution	Name in functions	Example function	Applications
Normal (Gaussian)	<code>norm</code>	<code>pnorm()</code>	Numerous (often assumed)
Student's $t$	<code>t</code>	<code>pt()</code>	Estimating parameters without knowing $\sigma$ ; $t$ -tests
Binomial (Bernoulli)	<code>binom</code>	<code>pbinom()</code>	Number of successes in <code>size</code> trials; logistic regression
Chi-squared ( $\chi^2$ )	<code>chisq</code>	<code>pchisq()</code>	Chi-squared tests (goodness of fit in a 2-way table)
$F$	<code>f</code>	<code>pf()</code>	ANOVA; $F$ -tests (model goodness-of-fit)

# Hypothesis tests

Hypothesis tests in base R tend to follow this format:

- a function call
- accepts vector or data frame inputs
- returns an object which can be assigned
- can be calculated “by hand” using other R functions
  - example: `t.test()` can also be found by calculating the  $t$  test statistic, followed by `pt()` to obtain a  $p$

# Student's $t$ -test

## One-sample $t$ -test

“Is sample mean  $\bar{x}$  different from population mean  $\mu$ ?”

$$t = \frac{\bar{x} - \mu_0}{s / \sqrt{n}}$$

```
1 mu <- 4147
2 x_bar <- mean(penguins$body_mass_g, na.rm=TRUE)
3 s <- sd(penguins$body_mass_g, na.rm=TRUE)
4 n <- penguins[!is.na(penguins$body_mass_g), "body_mass_g"]
5   pull() %>%
6   length()
7
8 t <- (x_bar - mu) / (s / sqrt(n))
9
10 # find p for this t score, on t distribution with
11 pt(q = t,
12    df = (n - 1),
13    lower.tail = TRUE)
14
15 # or:
16 t.test(x = penguins$body_mass_g,
17        mu = mu)
```

## Two-sample $t$ -test (independent)

“Is sample mean  $\bar{X}_1$  different from sample mean  $\bar{X}_2$ ?”

$$t = \frac{\bar{X}_1 - \bar{X}_2}{s_p \cdot \sqrt{\frac{1}{n_1} + \frac{1}{n_2}}}$$

```
1 adelie <- penguins %>%
2   filter(species == "Adelie")
3 chinstrap <- penguins %>%
4   filter(species == "Chinstrap")
5
6 t.test(x = adelie$body_mass_g,
7        y = chinstrap$body_mass_g)
```

# Chi-squared ( $\chi^2$ ) test

“Is the conditional distribution across two categorical variables equal, or varied?”

`chisq.test(x, y = NULL)` where `x` is a matrix, or `x` and `y` are vectors

`chisq_test(formula)` where `formula` is written in the format `response ~ explanatory`

```
1 # how are penguins distributed across species and islands?  
2 chisq_test(penguins, island ~ species)
```



⚠ This is not a good example of a chi-squared test, statistically speaking: our data do not meet the requirements for the test! But the example is left here because R will let us do it! (and it is syntactically correct)

# Tidy inference: **infer**

```
1 t <- penguins %>%
2   specify(response = body_mass_g) %>%
3   hypothesize(null = "point", mu = 4147) %>%
4   calculate(stat = "t") %>%
5   pull()
6
7 penguins %>%
8   specify(response = body_mass_g) %>%
9   assume(distribution = "t") %>%
10  visualize() +
11  shade_p_value(obs_stat = t, direction = "greater")
12
13 penguins %>%
14   t_test(response = body_mass_g,
15         mu = 4147)
16
17 penguins %>%
18   filter(species %in% c("Adelie", "Chinstrap")) %>%
19   t_test(formula = body_mass_g ~ species)
```



# Modeling

# Simple linear regression

“What is the relationship between  $X$  and  $Y$ ?”

Considering

$$Y = \beta_0 + \beta_1 X + \epsilon$$

where  $\beta_0$  is the  $y$  intercept,  $\beta_1$  is the slope coefficient,  $X$  is the explanatory or predictor variable,  $\epsilon$  is irreducible error, and  $Y$  is the response variable. We want to estimate (or “fit”)  $\hat{\beta}_0$  and  $\hat{\beta}_1$  such that the sum of squared residuals is minimized as much as possible.

`lm(formula, data)` fits a linear model and returns a model object, where formula is written `response ~ explanatory` and `data` is a data frame.

```
1 lm(flipper_length_mm ~ bill_length_mm,  
2   data = penguins) %>%  
3   summary()  
4  
5 lm_flipper_bill <- lm(flipper_length_mm ~ bill_length_mm, data = penguins)
```



# predict() outputs from a model

“For some  $x$ , what  $\hat{y}$  does the model predict?”

Models in R implement behavior for the `predict()` function, which supplies a data frame of input values  $X$  and returns a data frame containing corresponding  $Y$  values. The input data frame needs variable(s) of the same name(s) as the predictor(s). Or you can omit the data frame and get the fitted values.

💡 `predict()` functions are found in documentation under article names like `predict.lm`, even though the function call is `predict()`.

```
1 # min to max values, incrementing by 0.5:
2 x_range <- data.frame(bill_length_mm = seq(from=min(penguins$bill_length_mm, na.rm=TRUE), to=max(penguins$bi
3
4 predict(lm_flipper_bill, newdata = x_range)
5
6 # fitted values
7 predict(lm_flipper_bill)
```



# Tidy modeling: **parsnip**

```
1 linear_reg() %>%  
2   set_engine("lm") %>%  
3   fit(flipper_length_mm ~ bill_length_mm,  
4       data = penguins)  
5  
6 linear_reg() %>%  
7   set_engine("lm") %>%  
8   fit(flipper_length_mm ~ bill_length_mm,  
9       data = penguins)
```



# Model *responsibly*

- Consider the validity and reliability of your measures
- Correlation  $\neq$  causation
- Always keep in mind the assumptions made by your chosen method/model
- Validate unexpected results by recomputing “by hand” and/or running a different code implementation (or even in a different software)

# Wrap up

# Session in review

Today we learned about:

- how distributions and inference work in R
- linear modeling in R
- old and new ways of doing these things

Join us next week for machine learning intro!

